

$T(n) = T(n-1) + O(1) = O(n)$
 $T(n) = T(n-1) + O(n) = O(n^2)$
 $T(n) = T(n/2) + O(1) = O(\log n)$ (Binary search)
 $T(n) = T(n/2) + O(n) = O(n)$ (Quickselect)
 $T(n) = kT(n/k) + O(1) = O(n)$
 $T(n) = kT(n/k) + O(n) = O(n \log n)$, (Quicksort)
 $T(n) = T(n-1) + T(n-2) = O(2^n)$ (Fibonacci)
 $T(n) = T(n-1) + O(nk) = O(nk+1)$
 $T(n) = T(n-1) + O(\log n) = O(n \log n)$ (Stirling Approx)
 $T(n) = 2T(n-1) + 1 = O(2^n)$
 $T(n) = 2T(n/4) + O(1) = O(n^{0.5})$
 $n^b > (\log n)^a$ always
 $n^b < a^m$ always
 $O(n^{9.9}) < O(1.01^n)$

Binary search

```

int search(A, key, n)
begin = 0
end = n-1
while begin < end do:
mid = begin + (end-begin)/2;
if key <= A[mid] then
end = mid
else begin = mid+1
return (A[begin]==key) ? begin : -1
Invariant: A[begin] <= key <= A[end]
if key is in A.
T(n) = T(n/2) + O(1) = O(log n)

```

Peak finding

```

FindPeak(A, n)
if A[n/2] is a peak then return n/2
else if A[n/2+1] > A[n/2] then
Search for peak in right half.
else if A[n/2-1] > A[n/2] then
Search for peak in left half.
Invariant:
1. If we recurse in the right half,
then there exists a peak in the right half.
2. If we recurse in the right half,
then every peak in the right half is a peak
in the array.
T(n) = T(n/2) + O(1) = O(log n)

```

Sorting Algorithms	Time Complexities			Input Types			Stable	Advantages	Disadvantages	Invariant / Sorting Jumble Hax
	Best Case	Worst Case	Average Case	Best Case	Worst Case	Average Case				
Insertion Sort	$O(n)$	$O(n^2)$	$O(n^2)$	Already sorted/ almost sorted	Reverse sorted	Input chosen in random	Yes	Fast for almost sorted	Slow average case	After j iterations, first j elements sorted (not final), back part untouched
Selection Sort	$O(n^2)$	$O(n^2)$	$O(n^2)$	-	-	-	No	-	Slow for almost sorted array	After j iterations, first j elements correctly sorted in their final positions, original element swapped away
Bubble Sort	$O(n)$	$O(n^2)$	$O(n^2)$	Already sorted	Reverse sorted	Input chosen in random	Yes	Fast for almost sorted	Significantly slower for most inputs	After j iterations, last j elements correctly sorted in their final positions
Merge Sort	$O(n \log n)$	$O(n \log n)$	$O(n \log n)$	-	-	-	Yes	Fast on average	Slow for small n and sorted array Extra space required	After k iterations, the first k elements in the final array are sorted, each pair of elements are correctly sorted.
Quick Sort	$O(n \log n)$	$O(n^2)$	$O(n \log n)$	Random pivot	Every pivot is the smallest/biggest element	Random pivot	No	Optimizable, many partition and pivot choosing algorithms	Worst case is $O(n^2)$	After the partition step, every: arr[i] < pivot, where i < low arr[j] < pivot, where j > high Pivots always sorted

Quicksort Partition Algo $O(n)$:

```

partition(A[0..n], n, pIndex) // Assume no duplicates, n>1
pivot = A[pIndex]; // pIndex is the index of pivot
swap(A[0], A[pIndex]); // store pivot in A[0]
low = 1; // start after pivot in A[1]
high = n+1; // Define: A[n+1] = inf
while (low < high)
while (A[low] < pivot) and (low < high) do low++;
while (A[high] > pivot) and (low < high) do high--;
if (low < high) then swap(A[low], A[high]);
swap(A[1], A[low-1]); //swap the pivot back in
return low-1;

```

To deal with duplicates, 3-way partition algo.

BST

```

public TreeNode successor(){
if (rightTree != null)
return rightTree.searchMin();
TreeNode parent = parentTree;
TreeNode child = this;
while ((parent != null) && (child == parent.rightTree))
child = parent;
parent = child.parentTree;
}
return parent;
}

```

delete(v)

1. If v has two children, swap it with its successor.
2. Delete node v from binary tree (and reconnect children).
3. For every ancestor of the deleted node:
 - Check if it is height-balanced.
 - If not, perform a rotation.
 - Continue to the root.

Deletion may take up to $O(\log(n))$ rotations.

Things to note:

- BST may not be balanced
- Worst case runtime for Query Operations is $O(n)$
- For deletion of node v:
 - If no children, just delete v
 - If 1 child,
 - delete v + connect v.child to v.parent
 - If 2 children,
 - x = successor(v)
 - delete(x)
 - delete(v)
 - connect x to v.left, v.right, v.parent

Choice of Pivot:


- Pivot is good when it divides the array to 2 pieces, each size at least 1/10
- Median/First/Last: Worst case is the same $O(n^2)$
- Paranoid Quicksort:
 - Repeat the partition step until $p > n/10$ and $p < 9n/10$. Expected runtime is $O(n \log n)$ with high probability.

right-rotate(v) // assume v has left != null

```

w = v.left
w.parent = v.parent
v.parent = w
v.left = w.right
w.right = v

```



Quickselect kth smallest element

```

Select(A[1..n], n, k)
if (n == 1) then return A[1];
else choose random pivot index pIndex.
p = partition(A[1..n], n, pIndex)
if (k == p) then return A[p];
else if (k < p) then
return Select(A[1..p-1], k)
else if (k > p)
then return Select(A[p+1..n], k - p)
Invariant: The pivots will be in their final positions.

```

AVL trees: Height balanced if every node is height balanced

$|v.\text{left.height} - v.\text{right.height}| \leq 1$ and has at most $h < 2\log(n)$, and $n > 2^{h/2}$.

Insertion:

– Insert key in BST.

– Walk up tree:

- At every step, check for balance.
- If out-of-balance, use rotations to rebalance.

If v is out of balance and left heavy:

1. v.left is balanced: right-rotate(v)
2. v.left is left-heavy: right-rotate(v)
3. v.left is right-heavy: left-rotate(v.left), right-rotate(v)

If v is out of balance and right heavy:

1. v.right is balanced: left-rotate(v)
2. v.right is right-heavy: left-rotate(v)
3. v.right is left-heavy: right-rotate(v.right), left-rotate(v)

Trie vs AVL:

Time:

Trie tends to be faster: $O(L)$ vs $O(L_h)$.

Does not depend on size of total text.

Does not depend on number of strings.

Space:

Trie tends to use more space.

BST and Trie use $O(\text{text size})$ space.

But Trie has more nodes and more overhead.

Augmenting Data Structures

Most of the time, we will make use of existing

Data structures, but store additional

Data so that it works as we intend.

E.g. weight, rank, max, min

1. Maintain the Data Structure

(BST, array, linked list etc)

2. Modifying Operations

(Insert, delete)

3. Query Operations

Search, Select

Important!!

Make sure the additional data is updated during

Insertion/Deletion.

(a, b)-trees

B-trees (special type of (a,b)-trees where $a=B$ and $b = 2B$)

Search: check keylist. If key is not in keylist, need to look inside the relevant subtree.

If found, return node.

- An (a,b)-B-tree with n nodes has $O(\log_a n)$ height

- Binary search for a key at every node takes $O(\log_2 b)$ time

- Hence total search cost: $O(\log_2 b \cdot \log_a n)$ (Cost of searching a node \times Height)

$= O(\log_a n)$ since $\log_2 b$ is a constant $= O(\log n)$

Insertion: Traverse to the relevant node, then insert into keylist.

Rule 1: (a,b) child policy.

Rule 2: Internal nodes have 1 more child than its number of keys

Rule 3: All leaves have the same depth.

If keylist has $> b-1$ keys, perform split

1. Choose median key from keylist

2. Use that to split keylist into 2 halves.

3. Put median key to parent. Left half is left child, right half is right child.

- If parent is also over capacity, continue splitting upwards.

- To split root node, create a new node (the new root) to store the median.

Deletion:

1. Find successor of key to be deleted

2. Replace the key with its successor.

3. If node with less than $a-1$ keys after deletion, demote separating key from parent and join with sibling

Orthogonal Range Searching

Find split node:

Highest node where search includes both left and right subtree

Algorithm:

- $v = \text{FindSplit}(\text{low}, \text{high})$;

- $\text{LeftTraversal}(v, \text{low}, \text{high})$;

- $\text{RightTraversal}(v, \text{low}, \text{high})$;

$\text{FindSplit}(\text{low}, \text{high})$ $O(\log n)$

$v = \text{root}$;

$\text{done} = \text{false}$;

while !done {

if ($\text{high} \leq v.\text{key}$) then $v = v.\text{left}$;

else if ($\text{low} > v.\text{key}$) then $v = v.\text{right}$;

else ($\text{done} = \text{true}$);

}

return v ;

$\text{LeftTraversal}(v, \text{low}, \text{high})$

if ($\text{low} \leq v.\text{key}$) {

all-leaf-traversal($v.\text{right}$);

$\text{LeftTraversal}(v.\text{left}, \text{low}, \text{high})$;

}

else {

$\text{LeftTraversal}(v.\text{right}, \text{low}, \text{high})$;

}

Right is symmetric

Order statistics in BST:

Idea: Store the size of the subtree in each node

Select(k): select node of rank(k):

Time Complexity: $O(\log n)$

select(k)

rank = $m.\text{left}.\text{weight} + 1$;

if ($k == \text{rank}$) then

return v ;


else if ($k < \text{rank}$) then

return $m.\text{left}.\text{select}(k)$;

else if ($k > \text{rank}$) then

return $m.\text{right}.\text{select}(k - \text{rank})$;

Maintain weight during rotations:



Rank(v): return the rank of node v

Time Complexity: $O(\log n)$

rank(node)

rank = $\text{node}.\text{left}.\text{weight} + 1$;

while ($\text{node} \neq \text{null}$) do

if node is left child then

do nothing

else if node is right child then

rank += $\text{node}.\text{parent}.\text{left}.\text{weight} + 1$;

node = $\text{node}.\text{parent}$;

return rank;

Node type	#Keys		#Children	
	Min	Max	Min	Max
Root	1	$b - 1$	2	b
Internal	$a - 1$	$b - 1$	a	b
Leaf	$a - 1$	$b - 1$	0	0

More Algorithm Analysis

Sum of AP:

$d = \text{Common Difference}$

$a = \text{first term}$, then $a + (n - 1) \times d$ is the last term

$$S_n = \frac{n}{2}(2a + (n - 1) \times d) = \frac{n}{2}(\text{first term} + \text{last term})$$

Sum of GP:

$r = \text{Common Ratio}$

$a = \text{first term}$

$$S_n = \frac{a(r^n - 1)}{r - 1} = \frac{a(1 - r^n)}{1 - r}$$

If loop index $x =$ or \neq , there will be $\log_2 n$ number of steps

$$T(n) = kT\left(\frac{n}{k}\right) + O(n^2) = O(n^2)$$

$$T(n) = kT\left(\frac{n}{k}\right) + O(\sqrt{n}) = O(n)$$

$$T(n) = kT\left(\frac{n}{k}\right) + O(n \log n) = O(n \log n^2)$$

$$T(n) = kT\left(\frac{n}{k}\right) + O(\log n) = O(\log n^2)$$

Interval Trees:

Each node is an interval (a,b), sorted by left end-point (a).

Augment the maximum endpoint in subtree rooted at the node.

interval-search(x): find interval containing x ($O(\log n)$)

interval-search(x)

Augment: maximum endpoint in subtree

$c = \text{root}$;

while ($c \neq \text{null}$ and x is not in $c.\text{interval}$) do

if ($c.\text{left} == \text{null}$) then


$c = c.\text{right}$;

else if ($x > c.\text{left}.\text{max}$) then

$c = c.\text{right}$;

else $c = c.\text{left}$;

return $c.\text{interval}$;



if search goes right, then there is no interval in the left tree (key $>$ max of left subtree)

if search goes left, then there is no interval in the right tree (key $<$ left endpoint in left tree, so confirm $<$ anything in the right tree)

All-Overlaps Algorithm: ($O(k \log n)$)

Repeat until no more intervals:

- Search for interval.

- Add to list.

- Delete interval.

Repeat for all intervals on list:

- Add interval back to tree.

Chaining

Each bucket contains a linked list to store multiple

entries: Total space $= O(m+n)$, m = table size, n = linked list size.

Insert: compute $h(\text{key})$ $O(\text{cost}(h))$ and add (key, value) to linked list $O(1)$. Total: $O(1 + \text{cost}(h))$

Search: compute $h(\text{key})$ $O(\text{cost}(h))$ and look inside linked list $O(n)$ (worst case). Total: $O(n + \text{cost}(h))$

Simple Uniform Hashing Assumption:

• All keys equally likely to map to every bucket

• Keys are mapped independently

Expect number of entries per bucket $= n/m$.

Hence, expected search time $= O(1 + n/m) = O(1)$

Hashing

Symbol table has no easy way to sort: is not comparison based.

Direct Access Table: convert anything in the world into bits. Space required is $n^{\text{no. of bits}}$

Idea of Hashing: map n keys to m (usually less than n) buckets

Hash function: maps key to bucket $h: U \rightarrow \{1 \dots m\}$

Time Complexity: $O(\text{cost of finding } h(k) + \text{cost of accessing bucket})$

Collision: when $h(k_1) = h(k_2)$.

Unavoidable: table size $<$ universe size. By pigeonhole principle, collisions are bound to occur

Usually $O(1)$